


Multivariate Linear Regression

Housing Prices (Portland, OR)



(Single feature)

Size in feet ² (x)	Price (\$) in 1000's (y)
2104	460
1416	232
1534	315
852	178
...	...

(Single variate linear regression)

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

Predicting house price y
based on single feature of
house size x

Multiple features (variables).

Size (feet ²)	Number of bedrooms	Number of floors	Age of home (years)	Price (\$1000)
x_1	x_2	x_3	x_4	y
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178
...

Notataion

n : number of features

x^i : features of i -th training sample

x_j^i : value of feature j in i -th training sample

Hypothesis:


$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \dots + \theta_n x_n$$

$$x = \begin{bmatrix} x_0 \\ x_1 \\ \cdot \\ \cdot \\ \cdot \\ x_n \end{bmatrix} \in R^{n+1} \quad \theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \cdot \\ \cdot \\ \cdot \\ \theta_n \end{bmatrix} \in R^{n+1} \quad \theta_0: \text{bias, or} \\ \text{consider } x_0=1$$

In matrix form,

$$h_{\theta}(x) = x^T \theta$$

→ Multivariate linear regression



To represent h – need a structure/model & model parameters

Multivariate linear regression:

$$h_{\theta}(x) = x^T \theta$$

x : features of house

θ : parameter of model

Hypothesis $h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \dots + \theta_n x_n$

Parameters $\theta_0, \theta_1, \dots, \theta_n$

Cost function $J(\theta_0, \theta_1, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$

Update parameters $\theta_0 \theta_1 \dots \theta_n$ using gradient descent until convergence

Simultaneously update $\theta_j, j = 0, 1, \dots, n$, according to

$$\theta_j := \theta_j - \eta \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1, \dots, \theta_n)$$

η is a positive number, which is called the learning rate
– too small, slow learning; too large, divergence

Prepare Data

To make sure features are on a similar scale,
e.g., within $(-1, 1)$ range

Feature scaling

E.g. x_1 = size (0-2000 feet²)

x_2 = number of bedrooms (1-5)

$$x_1 = \frac{\text{size (feet}^2\text{)}}{2000}$$

$$x_2 = \frac{\text{number of bedrooms}}{5}$$

Mean normalization

Replace x_i with $x_i - \mu_i$ to make features have approximately zero mean
(Do not apply to $x_0 = 1$).

E.g. $x_1 = \frac{size - 1000}{2000}$

$$x_2 = \frac{\#bedrooms - 2}{5}$$

$$-0.5 \leq x_1 \leq 0.5, -0.5 \leq x_2 \leq 0.5$$

Use z score:

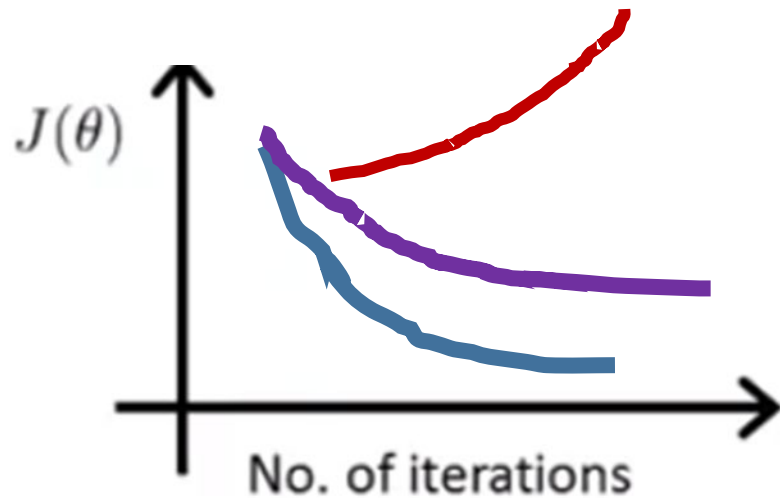


$$z = \frac{x - \mu}{\sigma}$$

where μ is the mean of feature x , and σ is the standard deviation of x .

If x is a normal distribution, z is a standard normal distribution (0 mean and unit variance)

Making sure gradient descent is working properly

Proper use of learning curve




Select a range of learning rate

$\eta = 0.001, 0.003, 0.01, 0.03, \dots$

Features and Polynomial Regression

(It is still linear regression but with polynomial features)

Polynomial Regression



Still consider a single physical feature: size (s).

Let :

$$x_1 = (\text{size}) \text{ or } x_1 = s$$

$$x_2 = (\text{size})^2 \text{ or } x_2 = s^2$$

$$x_3 = (\text{size})^3 \text{ or } x_3 = s^3$$


$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2$$

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3$$

Single variate polynomial regression as multi-variate linear regression

$$\begin{aligned} h_{\theta}(x) &= \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n \\ &= \theta_0 + \theta_1 s + \theta_2 s^2 + \cdots + \theta_n s^n \end{aligned}$$

Choice of features



$$h_{\theta}(\text{size}) = \theta_0 + \theta_1(\text{size}) + \theta_2(\text{size})^2$$

$$h_{\theta}(\text{size}) = \theta_0 + \theta_1(\text{size}) + \theta_2\sqrt{(\text{size})}$$

Solve for θ

- Gradient descent
- Normal equation (closed form solution)

How to optimize (minimize) a cost/loss function by tuning neural network weights for a predicted value to approach an actual value?

- Stochastic gradient descent (such as the LMS by Widrow and Hoff)
- Batch gradient descent as in linear regression, quadratic problem with least squares solution in closed form (may not be the best for large data set)
- Mini-batches (as in many deep networks)

m **examples** $(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})$; n **features**.

Examples: $m = 4$.

Size (feet ²)	Number of bedrooms	Number of floors	Age of home (years)	Price (\$1000)
x_1	x_2	x_3	x_4	y
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_4 x_4$$

$$X = \begin{bmatrix} 1 & 2104 & 5 & 1 & 45 \\ 1 & 1416 & 3 & 2 & 40 \\ 1 & 1534 & 3 & 2 & 30 \\ 1 & 852 & 2 & 1 & 36 \end{bmatrix} \quad y = \begin{bmatrix} 460 \\ 232 \\ 315 \\ 178 \end{bmatrix}$$

$$X\theta = y$$

$$\theta = (X^T X)^{-1} X^T y$$

m **training examples**, n **features**.

Gradient Descent

- Need to choose η
- Needs many iterations.

Normal Equation

- No need to choose η
- Don't need to iterate.

Use when n is large